

02 - RNN-LSTM Model Estimation

January 31, 2020

```
[1]: # If you run this it won't necessarily give you the best RNN model.  
# I've saved the best model that I estimated and I use it in the next Jupyter  
# Notebook.  
# This just an example of the steps I took for estimating the model.  
stock = 'Google'  
timesteps = 6 # This is the number of lags
```

0.1 Loading Data

```
[2]: import matplotlib.pyplot as plt  
import seaborn as sns  
sns.set()  
import pandas as pd
```

```
[3]: df = pd.read_csv(f'{stock}.csv')  
df.tail()
```

```
[3]:
```

	Date	Open	High	Low	Close \
3871	2020-01-06	1350.000000	1396.500000	1350.000000	1394.209961
3872	2020-01-07	1397.939941	1402.989990	1390.380005	1393.339966
3873	2020-01-08	1392.079956	1411.579956	1390.839966	1404.319946
3874	2020-01-09	1420.569946	1427.329956	1410.270020	1419.829956
3875	2020-01-10	1427.560059	1434.928955	1418.349976	1429.729980

	Adj Close	Volume
3871	1394.209961	1732300
3872	1393.339966	1502700
3873	1404.319946	1528000
3874	1419.829956	1500900
3875	1429.729980	1820700

```
[4]: data = pd.DataFrame()  
data[f'{stock}'] = df['Adj Close'].copy()
```

```
[5]: data[f'D{stock}'] = data[f'{stock}'].diff()  
data['Variance'] = (data[f'D{stock}']-data[f'D{stock}'].mean())**2
```

```
[6]: data.dropna(axis=0, inplace=True)
data.tail()
```

```
[6]:
```


	Google	DGoogle	Variance
3871	1394.209961	33.549927	1101.832552
3872	1393.339966	-0.869995	1.503220
3873	1404.319946	10.979980	112.867595
3874	1419.829956	15.510010	229.642085
3875	1429.729980	9.900024	91.087176

0.2 Splitting the data in training and test sample

```
[7]: data_training = pd.DataFrame()
data_test = pd.DataFrame()
data_training['Variance'] = data['Variance'][:-timesteps-3].copy()
data_test['Variance'] = data['Variance'][-timesteps-3:].copy()
```

```
[8]: data_training['Variance'].plot()
```

```
[8]: <matplotlib.axes._subplots.AxesSubplot at 0x2b9ef3e0908>
```



0.3 Training Data preprocessing

```
[9]: import numpy as np

x_train = []
y_train = []

for i in range(timesteps, data_training.shape[0]):
    x_train.append(data_training['Variance'].iloc[i-timesteps:i].values.
        →tolist())
    y_train.append((data_training['Variance'].iloc[i]))

x_train, y_train = np.array(x_train), np.array(y_train)

x_train.shape, y_train.shape
```

```
[9]: ((3860, 6), (3860,))
```

0.4 Scaling the data

```
[10]: # I am using logarithmic scale.
from sklearn.preprocessing import FunctionTransformer
scaler = FunctionTransformer(np.log1p, validate=True)
x_train = scaler.fit_transform(x_train)
x_train = x_train[..., np.newaxis]
x_train.shape
```

```
[10]: (3860, 6, 1)
```

0.5 Test data preprocessing

```
[11]: x_test = []
y_test = []

for i in range(timesteps, data_test.shape[0]):
    x_test.append(data_test['Variance'].iloc[i-timesteps:i].values.tolist())
    y_test.append((data_test['Variance'].iloc[i]))

x_test, y_test = np.array(x_test), np.array(y_test)

x_test.shape, y_test.shape
```

```
[11]: ((3, 6), (3,))
```

```
[12]: x_test = scaler.transform(x_test)
x_test = x_test[..., np.newaxis]
x_test.shape
```

```
[12]: (3, 6, 1)
```

0.6 Tensorflow Model

```
[13]: import tensorflow as tf
      from tensorflow.keras import Sequential
      from tensorflow.keras.layers import Dense, LSTM, GRU, Dropout,
      ↪BatchNormalization
      from tensorflow.keras.callbacks import ModelCheckpoint
```

```
[14]: model = Sequential()

      model.add(LSTM(units = 60, activation = 'relu', return_sequences = True,
      ↪input_shape = (timesteps, 1)))
      model.add(Dropout(0.2))
      model.add(BatchNormalization())

      model.add(LSTM(units = 60, activation = 'relu', return_sequences = True))
      model.add(Dropout(0.2))
      model.add(BatchNormalization())

      model.add(LSTM(units = 60, activation = 'relu'))
      model.add(Dropout(0.2))
      model.add(BatchNormalization())

      model.add(Dense(units = 1, activation = 'relu'))
```

```
[15]: opt = tf.keras.optimizers.Adam(lr=0.001, decay=1e-6)
```

```
[16]: model.compile(optimizer=opt, loss = 'mean_absolute_error')
```

```
[17]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 6, 60)	14880
dropout (Dropout)	(None, 6, 60)	0
batch_normalization (BatchNo	(None, 6, 60)	240
lstm_1 (LSTM)	(None, 6, 60)	29040
dropout_1 (Dropout)	(None, 6, 60)	0

batch_normalization_1 (Batch Normalization)	(None, 6, 60)	240

lstm_2 (LSTM)	(None, 60)	29040

dropout_2 (Dropout)	(None, 60)	0

batch_normalization_2 (Batch Normalization)	(None, 60)	240

dense (Dense)	(None, 1)	61
=====		
Total params: 73,741		
Trainable params: 73,381		
Non-trainable params: 360		

```
[18]: model.fit(x_train, y_train, epochs=50, batch_size=32, validation_split=0.0)
```

Train on 3860 samples

Epoch 1/50

3860/3860 [=====] - 13s 3ms/sample - loss: 79.5320

Epoch 2/50

3860/3860 [=====] - 6s 1ms/sample - loss: 78.2132

Epoch 3/50

3860/3860 [=====] - 6s 2ms/sample - loss: 77.4070

Epoch 4/50

3860/3860 [=====] - 5s 1ms/sample - loss: 77.0473

Epoch 5/50

3860/3860 [=====] - 5s 1ms/sample - loss: 76.5964

Epoch 6/50

3860/3860 [=====] - 5s 1ms/sample - loss: 76.6291

Epoch 7/50

3860/3860 [=====] - 5s 1ms/sample - loss: 76.5122

Epoch 8/50

3860/3860 [=====] - 4s 1ms/sample - loss: 76.4591

Epoch 9/50

3860/3860 [=====] - 6s 1ms/sample - loss: 76.3833

Epoch 10/50

3860/3860 [=====] - 6s 2ms/sample - loss: 76.3837

Epoch 11/50

3860/3860 [=====] - 5s 1ms/sample - loss: 76.4064

Epoch 12/50

3860/3860 [=====] - 5s 1ms/sample - loss: 76.2600

Epoch 13/50

3860/3860 [=====] - 4s 1ms/sample - loss: 76.1896

Epoch 14/50

3860/3860 [=====] - 5s 1ms/sample - loss: 76.1849

Epoch 15/50

3860/3860 [=====] - 4s 1ms/sample - loss: 76.2240

Epoch 16/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.1656
Epoch 17/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.0733
Epoch 18/50
3860/3860 [=====] - 5s 1ms/sample - loss: 75.9774
Epoch 19/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.3141
Epoch 20/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.0457
Epoch 21/50
3860/3860 [=====] - 6s 2ms/sample - loss: 76.0064
Epoch 22/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.0599
Epoch 23/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.0501
Epoch 24/50
3860/3860 [=====] - 5s 1ms/sample - loss: 75.9831
Epoch 25/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.3011
Epoch 26/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.0086
Epoch 27/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.0398
Epoch 28/50
3860/3860 [=====] - 5s 1ms/sample - loss: 75.7013
Epoch 29/50
3860/3860 [=====] - 5s 1ms/sample - loss: 75.9483
Epoch 30/50
3860/3860 [=====] - 5s 1ms/sample - loss: 75.8693
Epoch 31/50
3860/3860 [=====] - 4s 1ms/sample - loss: 76.1916
Epoch 32/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.1003
Epoch 33/50
3860/3860 [=====] - 5s 1ms/sample - loss: 75.7137
Epoch 34/50
3860/3860 [=====] - 5s 1ms/sample - loss: 75.8636
Epoch 35/50
3860/3860 [=====] - 5s 1ms/sample - loss: 75.9689
Epoch 36/50
3860/3860 [=====] - 5s 1ms/sample - loss: 75.9316
Epoch 37/50
3860/3860 [=====] - 5s 1ms/sample - loss: 75.4991
Epoch 38/50
3860/3860 [=====] - 5s 1ms/sample - loss: 76.0055
Epoch 39/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.7915

```

Epoch 40/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.7117
Epoch 41/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.9329
Epoch 42/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.6426
Epoch 43/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.9856
Epoch 44/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.8019
Epoch 45/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.5028
Epoch 46/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.4809
Epoch 47/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.6392
Epoch 48/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.8796
Epoch 49/50
3860/3860 [=====] - 4s 1ms/sample - loss: 75.6346
Epoch 50/50
3860/3860 [=====] - 6s 1ms/sample - loss: 75.3645

```

[18]: <tensorflow.python.keras.callbacks.History at 0x2ba6e2bc508>

0.7 Evaluating RNN-LSTM Model

```
[19]: y_pred = model.predict(x_test)
```

```
[20]: from sklearn.metrics import mean_absolute_error
mae = round(mean_absolute_error(y_test, y_pred), 3)
print(mae)
```

93.201

0.8 Visualization

```
[21]: plt.figure(figsize=(14,9))
plt.plot(y_test, color = 'black', label = 'Real Variance of Stock Price_
↳Returns')
plt.plot(y_pred, color = 'blue', label = 'RNN-LSTM Predicted Variance')
plt.title(f'Variance of Stock Price Returns Prediction for {stock}')
plt.xlabel('Date')
plt.ylabel('Variance')
plt.legend()
plt.show()
```

